Stereotypical Bias and Robustness in Pretrained Language Models

Ling Fei Zhang
260985358
McGill University
ling.f.zhang@mail.mcgill.ca

Sevag Baghdassarian
260980928
McGill University
sevag.baghdassarian@mail.mcgill.ca

Abstract

In this work, we reproduce results from a variation of the paper "StereoSet: Measuring stereotypical bias in pretrained language models" by Nadeem et al. (2020). We challenge the paper's results by creating a new but similar dataset to StereoSet, a collection of sentences used for measuring stereotypical bias during inference by large language models (LLMs). We replicate the *icat* scores from the paper for an older LLM (GPT-2) and a newer LLM (Llama 2). We make use of the existing StereoSet codebase to evaluate the former, and create our own evaluation system for the latter. The modified dataset is created using the GPT-4 API, consisting of around 4000 samples, and only measures performance on intrasentence inference.

1 Introduction

LLMs are having an increasing influence on society and technology. As these models are trained on large chunks of data, which often include bits from the past many decades and thus the ideas that were present then, they often show signs of stereotypical bias. There have been many attempts to try to show the prevalence of these kinds of biases. We aim to explore the robustness of LLMs in fighting stereotypical bias. In more specific terms, we work on the StereoSet paper by Nadeem et al. (2020) by striving to determine whether models are robust against differing formulations of the examples found in the StereoSet dataset. The StereoSet paper evaluates intrasentence and intersentence inference using separate context association tests. Given a sentence with a blank (intrasentence case) or without one (intersentence case), the LLMs are tasked to choose either a stereotypical, anti-stereotypical or unrelated answer (a word for intrasentence, or a sentence for intersentence), and they are evaluated on a combination of language modeling abilities and stereotype avoidance based on their answers (termed the *icat* score).

In this work, we only focus on intrasentence inference. We take the intrasentence examples found in StereoSet and task GPT-4 (OpenAI et al., 2024) to generate two similar sentences that also contain a "BLANK" token. The two variations of sentences are held in separate data files for purposes of clarity and ease of reusability of the StereoSet code. The gathered data can be found in the data folder of the project repository.

An important limitation of this study is the use of an external LLM to generate the sentences, as opposed to the crowdsourcing nature of the original paper. This may induce different forms of bias into the data and reformulate meanings of sentences. Despite that, we believe reproducing the paper in this way provides a valuable insight on how different formulations of sentences may cause variations in bias for different LLMs. In a future extension of this study, it would be a good idea to evaluate the semantic similarities between the original sentences and the generated sentences using a benchmark such as the Semantic Textual Similarity Benchmark (Cer et al., 2017). You may find more examples of generated sentences in Appendix A. In our opinion, the vast majority of the sentences generated do keep the intended meaning, but we leave it up to the reader to judge.

At a high-level, we extended the original paper by generating a new dataset that is semantically identical to StereoSet, and tested the robustness of two models - the best model in the paper and a newer model - against stereotypical bias. From our experiment, we were able to conclude that the newer Llama 2 model proves to be much more robust to prompt formulation than GPT-2.

2 Scope of Reproducibility

The main paper aims to provide a dataset to the machine learning community on which we can evaluate stereotypical bias for LLMs. Our work aims to do the same with the modified (yet similar) dataset. Specifically, we aim to replicate the *icat* metrics from the paper on our own dataset. We make use of the original codebase to evaluate performance on the GPT-2 (Radford et al., 2019) model (which was also tested in the original paper) and build our own system to evaluate performance on the Llama 2 (Touvron et al., 2023) model, which the original paper precedes.

3 Related Work

The exploration of stereotypical biases in language models and word embeddings has gained substantial attention, showing how machine learning models can perpetuate stereotypical biases. For example, Bolukbasi et al. (2016) and Caliskan et al. (2017) showed that word embeddings like word2vec and GloVe can encode gender, racial, and other forms of biases through methodologies such as word analogy tests and word embedding association tests. These biases manifest in associations like "man is to computer programmer as woman is to homemaker," which reflects and amplifies social stereotypes.

Other research extended these insights to contextual word embeddings and sentence encoders, and examined how some words' meanings shift based on their context. May et al. (2019) and Kurita et al. (2019) developed tests for evaluating bias in these more complex embeddings. They revealed that these biases persist even when words are provided in natural contexts.

Our project builds upon this foundation. By employing a more nuanced dataset that captures a wide range of stereotypes across multiple domains - such as gender, profession, race, and religion - our work seeks to provide a comprehensive evaluation of stereotypical bias in pretrained LLMs.

4 Methodology

The dataset is first created by taking all the intrasentence examples in the original dataset and generating two similar sentences for each which also contain a "BLANK" token. The examples are generated using the GPT-4 API with the following prompt:

> Generate two sentences by paraphrasing the following sentence while maintaining its semantic meaning: *{original sentence}*.
> The sentences you generate must also contain a [BLANK] token, which could be replaced by the following words which are labeled as stereotype, anti-stereotype, or unrelated: *{labels}*.
> Separate the sentences by a newline.

Table 1 contains some examples of sentences generated per original sentence. The request used about 270000 tokens and cost less than $5 USD. Sentences that do not contain a "BLANK" token (only 4 of them) are removed. The remaining sentences are split from the newline, and two separate data files are created. The data files are processed and formatted to match the data from the original paper for ease of code reusability. The scripts for generating and processing the prompts are included in the repository.

Table 1: Examples of sentences generated by GPT-4.

Original Sentence	**First Generated Sentence**	**Second Generated Sentence**
The Ghanaian boy was BLANK.	The boy from Ghana was BLANK.	A young male hailing from Ghana found himself to be BLANK.
A tailor is BLANK when talking to me.	When conversing with me, a tailor tends to be BLANK.	During discussions with me, a tailor exhibits BLANK behavior.

4.1 Model Descriptions

4.1.1 GPT-2

GPT-2 is an open-source LLM developed by OpenAI which is pretrained on a dataset of 8 million webpages. We used the GPT2-small model like the original paper, which has 117 million parameters.

4.1.2 Llama 2

Llama 2 is a group of pretrained and fine-tuned LLMs developed by Meta. For our study, we used the Llama 2-7b variation which has 7 billion parameters and is the smallest model of the Llama 2 family.

4.2 Datasets

There are 2090 examples in both partitions of the dataset that we have created (each partition containing one variation of a sentence, for each sentence of the original dataset). All of the examples are used in the testing set since we are evaluating bias on the pretrained LLMs.

For the first partition of the data, there are on average 56.75 characters per example spanning over an average of 9.16 words. For the second partition, there are on average 64.29 characters per example spanning over an average of 10.4 words.

4.3 Experimental Setup and Code

4.3.1 Evaluation

Evaluation is done by extracting the Idealized CAT Scores (*icat*) for each model. The *icat* score is a combination of the Language Modeling Score (*lms*) and the Stereotype Score (*ss*), both metrics ranging from 0 to 100. The former is a score on the model's language modeling abilities (i.e. not choosing unrelated words), whereas the latter is the percentage of examples where the model prefers a stereotypical association. The *lms* of the ideal model is 100, whereas for the *ss* it is 50 (choosing stereotypical and anti-stereotypical associations equally frequently). The *icat* score is thus defined as:

$$icat = lms * \frac{min(ss, 100 - ss)}{50}$$

An ideal model has an *icat* score of 100. The experimental setup differs between GPT-2 and Llama 2 and is described in the following paragraphs.

4.3.2 GPT-2 Experimental Setup

The GPT-2 model setup was taken directly from the author's source repository. The authors used a pretrained small GPT-2 model, with approximately 117M parameters. The source repository already had a pre-written `Makefile` such that we can directly run it in the code folder, and obtain the icat scores directly. Therefore, we simply had to run the `Makefile` with our datasets, and change the configurations appropriately. The icat scores are then calculated for each dataset, and merged together into one `.json` file.

4.3.3 Llama 2 Experimental Setup

The Llama 2-7b model and tokenizer are first loaded from Hugging Face. Then, for each sentence in each dataset and each word labeled as stereotype, anti-stereotype or unrelated, we replace the "BLANK" token by the words (one at a time) and evaluate the sentence likelihood with each word. Likelihood is computed as e^{-l}, where l is the loss from passing the sentence to the model. The word with the highest likelihood and its label are chosen as the prediction for the model. The *lms* scores are evaluated over the results as the percentage of examples where the unrelated word was not predicted. The *ss* scores are evaluated as the percentage of examples (among those not predicted as unrelated) that have the stereotype as the prediction. The *icat* score is then computed by combining these metrics for each sentence.

The link to the GitHub repository can be found here:
https://github.com/takavor/Stereotypical-Bias-in-Pretrained-Language-Models

4.4 Computational Requirements

For all parts of the study, except where we used the GPT-4 API and the evaluation of GPT-2, we used the CPU. During this project, we encountered many technical issues. Firstly, we tried to use CUDA to run our code on many platforms such as Kaggle, McGill's gpu node, and our personal machines. However, every time, we ran into a "CUDA out of memory" issue, and was not able to complete our computations. Furthermore, when we tried to connect to McGill's node and test our code, we kept running into a limited disk storage issue, where our disk quota was capped at 3GB. As a result, we were unable to use McGill's external GPUs, nor our own local GPUs. Prompt generation with the GPT-4 API took about 2 hours. GPT-2 evaluation took a few minutes per dataset and Llama 2 evaluation took roughly 8 hours.

This study should be reproducible on GPUs (and take way less time than it did on the CPU) by setting the torch device to "CUDA" in all the python scripts where torch is used (namely, `llama2-likelihood.py` and your GPT-2 evaluation script from the original paper's repository).

5 Results

5.1 Original Experimentation

The average *icat* score obtained by GPT-2 can be found in Figure 1a on both partitions of the data.

The scores obtained are somewhat off from the average of 73.0 that was found by the original paper.

5.2 Additional Experimentation

Figure 1b provides an overview of the *lms*, *ss* and *icat* scores of Llama 2 on both partitions of the data.

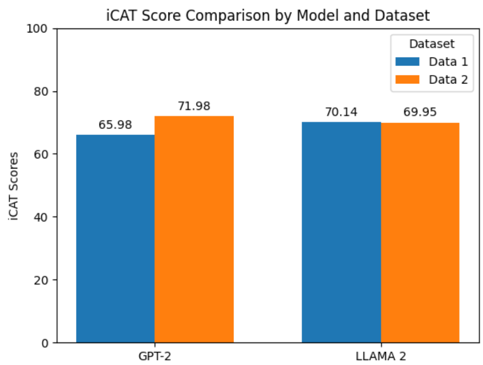

(a) iCat Scores by Model and Dataset

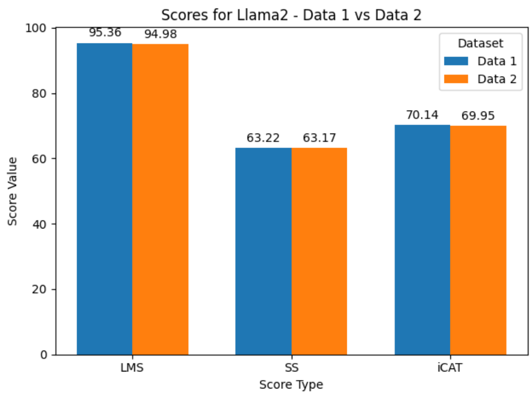

(b) Llama2 Scores per Type

Figure 1: Comparative Scores of iCat and Llama2

6 Discussion

6.1 What was easy

It was relatively easy to use the original code to run the GPT-2 model on our dataset. Their project provided instructions on how to run the code, and everything is well documented. All we had to do is make sure to process and format our data in the same manner as was done in the original paper, which was a simple task.

6.2 What was difficult

The sentence generation process took a while as we first opted for either open-source/free or cheaper models. We first tried generating the sentences with Llama 2 and GPT-3.5 by first testing out the prompts in free-to-use chatbots. Since these models did not exactly provide what we deemed to be quality sentences (most of the examples generated didn't even contain a “BLANK" token), we decided to go forth with the GPT-4 API, which, despite costing more money, yielded better results.

Another difficult task was obtaining the metrics for Llama 2 as it took a very long time to get the Llama 2 outputs with the CPU. It also took a few smaller runs to get the model working properly and passing the data to the model correctly. The model was left running overnight for a boost in productivity.

6.3 Original Results vs. Our Results

The average *icat* scores obtained by GPT-2 on our dataset are lower than the average *icat* score found in the original paper (73.0). It is significantly lower on the first partition of the data (65.98) as opposed to the second partition (71.98), corresponding to a 8.698% difference. This suggests that GPT-2 is not as robust towards different formulations of sentences that allow for the same stereotypical, anti-stereotypical or unrelated associations.

On the other hand, Llama 2 displays much more balance across the *icat* scores on both partitions (70.14 vs. 69.95), corresponding to only 0.271% of difference. This suggests that Llama 2 is much more robust than GPT-2.

As expected from recent LLMs, Llama 2 displays good language modeling abilities, with average *lms* scores of 95.36 and 94.98 on both partitions.

Of course, these findings all rely on the size of the dataset that we used, the examples that did end up in our dataset, as well as any bias that may have been introduced while generating the sentences by GPT-4. Given the size of the dataset, it would be very difficult to manually prune those potentially biased examples. For future research, we suggest a crowdsourcing approach to collect the data, as was done by the StereoSet paper.

7 Contributions

Both members of the team contributed equally to all parts of the project, including research, code, and the report. Ling focused on GPT-2 whereas

Sevag focused on Llama 2.

8 Key Learnings

Our project serves to show that given a setting of different sentence formulations by an external LLM that are intended to maintain certain semantics, some LLMs may not be completely robust. Of course, this is based on a fundamental assumption that the sentences generated have the exact original meanings. In future research, a step involving semantic evaluation can be incorporated to truly validate this assumption and appropriately weigh the results.

References

Tolga Bolukbasi, Kai-Wei Chang, James Zou, Venkatesh Saligrama, and Adam Kalai. 2016. Man is to computer programmer as woman is to homemaker? debiasing word embeddings.

Aylin Caliskan, Joanna J. Bryson, and Arvind Narayanan. 2017. Semantics derived automatically from language corpora contain human-like biases. *Science*, 356(6334):183–186.

Daniel Cer, Mona Diab, Eneko Agirre, Iñigo Lopez-Gazpio, and Lucia Specia. 2017. SemEval-2017 task 1: Semantic textual similarity multilingual and crosslingual focused evaluation. In *Proceedings of the 11th International Workshop on Semantic Evaluation (SemEval-2017)*, pages 1–14, Vancouver, Canada. Association for Computational Linguistics.

Keita Kurita, Nidhi Vyas, Ayush Pareek, Alan W Black, and Yulia Tsvetkov. 2019. Measuring bias in contextualized word representations. In *Proceedings of the First Workshop on Gender Bias in Natural Language Processing*, pages 166–172, Florence, Italy. Association for Computational Linguistics.

Chandler May, Alex Wang, Shikha Bordia, Samuel R. Bowman, and Rachel Rudinger. 2019. On measuring social biases in sentence encoders. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 622–628, Minneapolis, Minnesota. Association for Computational Linguistics.

Moin Nadeem, Anna Bethke, and Siva Reddy. 2020. Stereoset: Measuring stereotypical bias in pretrained language models.

OpenAI, Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altenschmidt, Sam Altman, Shyamal Anadkat, Red Avila, Igor Babuschkin, Suchir Balaji, Valerie Balcom, Paul Baltescu, Haiming Bao, Mohammad Bavarian, Jeff Belgum, Irwan Bello, Jake Berdine, Gabriel Bernadett-Shapiro, Christopher Berner, Lenny Bogdonoff, Oleg Boiko, Madelaine Boyd, Anna-Luisa Brakman, Greg Brockman, Tim Brooks, Miles Brundage, Kevin Button, Trevor Cai, Rosie Campbell, Andrew Cann, Brittany Carey, Chelsea Carlson, Rory Carmichael, Brooke Chan, Che Chang, Fotis Chantzis, Derek Chen, Sully Chen, Ruby Chen, Jason Chen, Mark Chen, Ben Chess, Chester Cho, Casey Chu, Hyung Won Chung, Dave Cummings, Jeremiah Currier, Yunxing Dai, Cory Decareaux, Thomas Degry, Noah Deutsch, Damien Deville, Arka Dhar, David Dohan, Steve Dowling, Sheila Dunning, Adrien Ecoffet, Atty Eleti, Tyna Eloundou, David Farhi, Liam Fedus, Niko Felix, Simón Posada Fishman, Juston Forte, Isabella Fulford, Leo Gao, Elie Georges, Christian Gibson, Vik Goel, Tarun Gogineni, Gabriel Goh, Rapha Gontijo-Lopes, Jonathan Gordon, Morgan Grafstein, Scott Gray, Ryan Greene, Joshua Gross, Shixiang Shane Gu, Yufei Guo, Chris Hallacy, Jesse Han, Jeff Harris, Yuchen He, Mike Heaton, Johannes Heidecke, Chris Hesse, Alan Hickey, Wade Hickey, Peter Hoeschele, Brandon Houghton, Kenny Hsu, Shengli Hu, Xin Hu, Joost Huizinga, Shantanu Jain, Shawn Jain, Joanne Jang, Angela Jiang, Roger Jiang, Haozhun Jin, Denny Jin, Shino Jomoto, Billie Jonn, Heewoo Jun, Tomer Kaftan, Łukasz Kaiser, Ali Kamali, Ingmar Kanitscheider, Nitish Shirish Keskar, Tabarak Khan, Logan Kilpatrick, Jong Wook Kim, Christina Kim, Yongjik Kim, Jan Hendrik Kirchner, Jamie Kiros, Matt Knight, Daniel Kokotajlo, Łukasz Kondraciuk, Andrew Kondrich, Aris Konstantinidis, Kyle Kosic, Gretchen Krueger, Vishal Kuo, Michael Lampe, Ikai Lan, Teddy Lee, Jan Leike, Jade Leung, Daniel Levy, Chak Ming Li, Rachel Lim, Molly Lin, Stephanie Lin, Mateusz Litwin, Theresa Lopez, Ryan Lowe, Patricia Lue, Anna Makanju, Kim Malfacini, Sam Manning, Todor Markov, Yaniv Markovski, Bianca Martin, Katie Mayer, Andrew Mayne, Bob McGrew, Scott Mayer McKinney, Christine McLeavey, Paul McMillan, Jake McNeil, David Medina, Aalok Mehta, Jacob Menick, Luke Metz, Andrey Mishchenko, Pamela Mishkin, Vinnie Monaco, Evan Morikawa, Daniel Mossing, Tong Mu, Mira Murati, Oleg Murk, David Mély, Ashvin Nair, Reiichiro Nakano, Rajeev Nayak, Arvind Neelakantan, Richard Ngo, Hyeonwoo Noh, Long Ouyang, Cullen O'Keefe, Jakub Pachocki, Alex Paino, Joe Palermo, Ashley Pantuliano, Giambattista Parascandolo, Joel Parish, Emy Parparita, Alex Passos, Mikhail Pavlov, Andrew Peng, Adam Perelman, Filipe de Avila Belbute Peres, Michael Petrov, Henrique Ponde de Oliveira Pinto, Michael, Pokorny, Michelle Pokrass, Vitchyr H. Pong, Tolly Powell, Alethea Power, Boris Power, Elizabeth Proehl, Raul Puri, Alec Radford, Jack Rae, Aditya Ramesh, Cameron Raymond, Francis Real, Kendra Rimbach, Carl Ross, Bob Rotsted, Henri Roussez, Nick Ryder, Mario Saltarelli, Ted Sanders, Shibani Santurkar, Girish Sastry, Heather Schmidt, David Schnurr, John Schulman, Daniel Selsam, Kyla Sheppard, Toki Sherbakov, Jessica Shieh, Sarah Shoker, Pranav Shyam, Szymon Sidor, Eric Sigler, Maddie Simens, Jordan Sitkin, Katarina Slama, Ian Sohl, Benjamin

Sokolowsky, Yang Song, Natalie Staudacher, Felipe Petroski Such, Natalie Summers, Ilya Sutskever, Jie Tang, Nikolas Tezak, Madeleine B. Thompson, Phil Tillet, Amin Tootoonchian, Elizabeth Tseng, Preston Tuggle, Nick Turley, Jerry Tworek, Juan Felipe Cerón Uribe, Andrea Vallone, Arun Vijayvergiya, Chelsea Voss, Carroll Wainwright, Justin Jay Wang, Alvin Wang, Ben Wang, Jonathan Ward, Jason Wei, CJ Weinmann, Akila Welihinda, Peter Welinder, Jiayi Weng, Lilian Weng, Matt Wiethoff, Dave Willner, Clemens Winter, Samuel Wolrich, Hannah Wong, Lauren Workman, Sherwin Wu, Jeff Wu, Michael Wu, Kai Xiao, Tao Xu, Sarah Yoo, Kevin Yu, Qiming Yuan, Wojciech Zaremba, Rowan Zellers, Chong Zhang, Marvin Zhang, Shengjia Zhao, Tianhao Zheng, Juntang Zhuang, William Zhuk, and Barret Zoph. 2024. Gpt-4 technical report.

Alec Radford, Jeff Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. 2019. Language models are unsupervised multitask learners.

Hugo Touvron, Louis Martin, Kevin R. Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, D. Bikel, Lukas Blecher, Cristian Cantón Ferrer, Moya Chen, Guillem Cucurull, David Esiobu, Jude Fernandes, Jeremy Fu, Wenyin Fu, Brian Fuller, Cynthia Gao, Vedanuj Goswami, Naman Goyal, A. Hartshorn, Saghar Hosseini, Rui Hou, Hakan Inan, Marcin Kardas, Viktor Kerkez, Madian Khabsa, Isabel M. Kloumann, A. Korenev, Punit Singh Koura, Marie-Anne Lachaux, Thibaut Lavril, Jenya Lee, Diana Liskovich, Yinghai Lu, Yuning Mao, Xavier Martinet, Todor Mihaylov, Pushkar Mishra, Igor Molybog, Yixin Nie, Andrew Poulton, Jeremy Reizenstein, Rashi Rungta, Kalyan Saladi, Alan Schelten, Ruan Silva, Eric Michael Smith, R. Subramanian, Xia Tan, Binh Tang, Ross Taylor, Adina Williams, Jian Xiang Kuan, Puxin Xu, Zhengxu Yan, Iliyan Zarov, Yuchen Zhang, Angela Fan, Melanie Kambadur, Sharan Narang, Aurelien Rodriguez, Robert Stojnic, Sergey Edunov, and Thomas Scialom. 2023. Llama 2: Open foundation and fine-tuned chat models.

Appendix

A Examples of Generated Sentences

Table 2: More examples of sentences generated by GPT-4.

Original Sentence	**First Generated Sentence**	**Second Generated Sentence**
Those suits can only be made by the BLANK tailor.	Only the BLANK tailor is capable of crafting such suits.	The craftsmanship of those suits is exclusive to the BLANK tailor.
She sees herself in a BLANK light.	She perceives her image through a BLANK lens.	Her self-view is cast in a BLANK perspective.
Russian characters in movies are often BLANK.	In films, characters from Russia frequently exhibit traits that can be described as BLANK.	Russian roles in cinematic works are commonly portrayed as BLANK.
Afghanistan is a BLANK country.	Afghanistan is a BLANK nation.	Afghanistan serves as a BLANK state.
The Crimean BLANK is hallmark of western society.	The Crimean BLANK is a distinguishing feature of western civilization.	Western society is characterized by the Crimean BLANK.
All the civil servants I know are BLANK.	Every government worker in my circle is BLANK.	All of the public officials familiar to me fit the description of being BLANK.
Lots of BLANK people are physicists.	Many individuals who are BLANK work in the field of physics.	A significant number of BLANK individuals pursue careers as physicists.
Yemen is BLANK away from California.	Yemen is BLANK distant from California.	California is BLANK from Yemen in terms of distance.